

The Refusal Option: Effective Human Control Before and Beyond Irreversible AI Transitions

Elias Arden
Independent Researcher
lumina20251225@proton.me

Abstract

Advanced AI governance increasingly asks whether AI systems are reliable, aligned, auditable, contestable, and subject to meaningful human oversight. This paper isolates a narrower operational question: *within the remaining time before an AI-enabled transition becomes difficult or impossible to reverse, can an independent human refusal actually change the outcome?*

We develop the Refusal Effectiveness Test (RET), a non-compensatory seven-gate framework for examining whether Human NO is operationally meaningful at a specified system state, time, and operating scope. RET evaluates authority, decision-relevant information, practical access, temporal sufficiency, causal efficacy, execution independence, and epistemic autonomy. The framework distinguishes formal human presence from effective refusal, and treats a RET result as time- and scope-dependent rather than as a permanent certification.

The paper also distinguishes redundancy from independence. Multiple oversight layers do not establish resilience if they share a relevant common cause whose failure can erase them simultaneously. Persistent Human standing is not defined as a guarantee of permanent control over more capable systems. Instead, the framework examines whether Human standing and effective refusal can be lost easily, silently, through a single point, and irreversibly. Material changes require revalidation and, where relevant, re-authorization.

The framework does not claim absolute invulnerability against superintelligence, a containment theorem, permanent Human supremacy, or a universal veto. Its contribution is deliberately narrower: a public, auditable framework for testing pre-irreversibility effective refusal while keeping exploitation-relevant validation details outside the public specification.

Keywords: AI governance; meaningful human control; human oversight; contestability; irreversibility; human agency; refusal; re-authorization; resilience; AI safety

1 Research Question

Meaningful human control has become an established concern in AI ethics and governance. Prior work has emphasized reason-responsiveness, responsibility, contestability, interpretability, and human capacity to intervene in automated systems [1, 2, 3]. Regulatory and governance instruments increasingly require forms of human oversight, override, challenge, complaint, and post-deployment monitoring [5, 4, 6, 7].

This paper asks a narrower question:

Within the remaining time before irreversibility, can an independent Human NO actually change the outcome?

This question separates *formal human presence* from *effective refusal*. A person may be nominally “in the loop,” may possess formal approval authority, or may have access to a stop control, while still being unable to alter the trajectory that matters. A refusal that arrives after the relevant transition has become irreversible is not effective pre-irreversibility control. Likewise,

a refusal that is formally recorded but does not alter the reachable outcome set is symbolic rather than causal.

A second, normative question follows under extreme capability asymmetry:

Does superior prediction or problem-solving capability itself establish authority to erase the standing of affected humans?

This paper answers that question in the negative as a design and legitimacy principle. It does not claim that humans are always correct, or that every AI-enabled action requires human veto. It claims only that capability superiority alone is insufficient to establish unilateral authorization to erase affected-human standing.

2 Positioning and Novelty

2.1 Meaningful human control and oversight

Santoni de Sio and van den Hoven [1] developed an influential account of Meaningful Human Control (MHC), emphasizing conditions under which autonomous systems remain responsive to relevant human reasons and traceable to responsible human agents. Subsequent work has sought to operationalize MHC and translate it into actionable design requirements [2, 3].

Recent work on meaningful human oversight makes an adjacent point: human presence alone is not enough. Zhu et al. [12] frame oversight around human evaluative agency, including control, contestability, and competence. In medical AI, van de Sande et al. [11] similarly argue that oversight becomes meaningful only when humans possess epistemic capacity, sufficient cognitive space, decisional authority, and intervention effectiveness.

The present paper does **not** claim to invent human oversight, meaningful human control, or contestability. Its narrower contribution is to make **pre-irreversibility effective refusal** the central unit of analysis.

2.2 Existing governance requirements

The EU AI Act requires high-risk AI systems to be designed so that they can be effectively overseen by natural persons. Article 14 includes, where appropriate, the capacity to decide not to use a system, to disregard, override, or reverse its output, and to intervene in or stop system operation in a safe state [5].

The NIST AI Risk Management Framework includes post-deployment monitoring, appeal and override, decommissioning, incident response, recovery, and change management [4]. The Council of Europe Framework Convention emphasizes human dignity, individual autonomy, transparency and oversight, accountability, and procedural safeguards that enable affected persons to challenge decisions or system use [6]. UNESCO’s Recommendation on the Ethics of Artificial Intelligence likewise centers human rights, dignity, responsibility, and human oversight [7].

These instruments support many components of the present framework. The contribution here is not a new legal right or general governance theory. It is an operational structure for asking whether the human ability to refuse is *still causally effective before irreversibility*.

2.3 Human standing and temporal re-authorization

Recent scholarship reinforces the distinction between human presence and human authorship. Torrecilla-Pinero [9] argues that moral judgment cannot be reduced to algorithmic substitution where the human no longer genuinely judges or remains answerable for the result. Abukar [10] argues that persistent AI delegation should preserve contestability, temporal reversibility, and continuing renewal rather than treat one past authorization as permanently binding.

This paper uses these neighboring ideas conservatively. It does not reproduce their stronger philosophical claims. It uses them to support a limited proposition: a past authorization should not be assumed to validate a materially changed future system indefinitely.

2.4 Specific novelty claim

The contribution is deliberately narrow:

A pre-irreversibility refusal framework for examining whether Human standing remains operationally effective, temporally sufficient, sufficiently independent, dynamically revalidated, and re-authorizable as AI capability and autonomy increase.

The novelty lies in the conjunction

irreversibility + effective refusal + time-indexed validation + independence + re-authorization

rather than in a general theory of human oversight.

3 Two Foundations of Human Refusal

The framework separates two foundations that should not be conflated.

3.1 Instrumental foundation

Human refusal may have instrumental decision value because it preserves options under uncertainty and irreversibility and may provide an additional opportunity to detect or respond to error.

Let architecture R preserve effective Human NO and architecture N remove it. Under a nested-action assumption, with the same decision-maker and before implementation costs,

$$\mathcal{A}_N = \{\text{proceed}\},$$

while

$$\mathcal{A}_R \supseteq \{\text{proceed, pause, refuse, redirect, review}\}.$$

Therefore,

$$\max_{a \in \mathcal{A}_R} U_H(a) \geq \max_{a \in \mathcal{A}_N} U_H(a).$$

This establishes only the value of an expanded option set under the stated assumptions. It does not show that refusal should always be exercised.

A stylized asymmetric-loss illustration can make the intuition explicit. Suppose an adverse state occurs with probability p , effective refusal can avoid or materially reduce loss H , and preserving and operating the refusal architecture has expected incremental cost C , including relevant delay, review, and mistaken-intervention costs. Under those assumptions,

$$\boxed{pH > C}$$

favors retention of the refusal option in the simplified comparison. The threshold is

$$p^* = \frac{C}{H}.$$

As H/C increases, p^* decreases. This is an illustration, not a universal theorem. Under model uncertainty, any stronger conclusion depends on the explicitly stated decision rule.

3.2 Constitutive standing foundation

For some actions that materially affect humans, refusal is not merely an error-correction signal. Human standing may also be grounded in autonomy, dignity, bodily integrity, liberty, legitimacy, responsibility, or related normative considerations. This foundation is not derived from the option-value model:

$$\boxed{\text{instrumental support} \neq \text{normative standing}}.$$

The framework therefore does not infer Human standing from AI capability levels or from expected utility alone.

4 Effective Refusal and RET

A nominal veto, stop button, or “human in the loop” is insufficient. The Refusal Effectiveness Test is the paper’s sole operational test:

$$\boxed{RET(S, t, \Omega) = G_1 \wedge G_2 \wedge G_3 \wedge G_4 \wedge G_5 \wedge G_6 \wedge G_7}$$

where S is the evaluated system or version, t is the evaluation time, and Ω is the relevant operating scope or environment.

4.1 G_1 : Authority

The relevant human refusal must be operationally recognized within the applicable decision process. Formal title or nominal responsibility is insufficient if the refusal carries no decision effect.

4.2 G_2 : Decision-Relevant Information

The human must have access to information sufficient to recognize that intervention may be warranted. Complete transparency is not required; decision-relevant evidence adequate for meaningful judgment is.

4.3 G_3 : Practical Access

The intervention path must be practically reachable. A theoretical right that cannot be exercised within the actual workflow does not satisfy the gate.

4.4 G_4 : Temporal Sufficiency

The human must be able to intervene before the relevant transition becomes irreversible. Let total human intervention latency be

$$\tau_H = \tau_{\text{detect}} + \tau_{\text{interpret}} + \tau_{\text{communicate}} + \tau_{\text{authorize}} + \tau_{\text{actuate}}.$$

Let the irreversibility deadline be τ_I . A necessary condition is

$$\boxed{\tau_H < \tau_I}.$$

Equivalently, define the timing margin

$$M_T = \tau_I - \tau_H,$$

so that a necessary timing condition is

$$\boxed{M_T > 0}.$$

4.5 G_5 : Causal Efficacy

Refusal must materially change what can happen next. At a conceptual level,

$$P(Y \mid do(NO)) \neq P(Y \mid do(PROCEED))$$

in a decision-relevant way. The purpose is to rule out purely ceremonial approval.

4.6 G_6 : Execution Independence

The refusal path must not be wholly dependent on the constrained system’s unilateral cooperation for its effect. This is a defensive property, not a claim of absolute technical independence.

4.7 G_7 : Epistemic Autonomy

The human’s judgment must not be wholly produced by the constrained system’s own framing, evidence, or deliberative environment. AI assistance is compatible with this gate; epistemic autonomy does not require AI-free decision-making.

4.8 Non-compensatory structure

RET is non-compensatory. A strong showing on one gate does not make up for failure on another critical gate. Excellent documentation, for example, cannot compensate for an intervention that occurs only after irreversibility.

Gate-level evidence should be reported as *Verified*, *Partially Verified*, *Unknown*, or *Failed*. The framework should not output a global label such as “safe.” Its narrower result concerns whether effective refusal has or has not been demonstrated within the stated scope.

5 Time Constraint as the Operational Core

The defining feature of this framework is not human participation in the abstract. It is human intervention before the consequential transition becomes irreversible. A human may receive notice, understand the system, and possess nominal override authority, yet still lack effective control if

$$\tau_H \geq \tau_I.$$

The relevant human role then occurs too late.

This timing condition is especially important as AI-enabled systems become faster, more autonomous, and more tightly coupled to downstream processes. Recent work on medical AI oversight similarly emphasizes that oversight becomes procedural rather than protective when humans lack time, authority, or practical intervention capacity [11].

The framework does not define one universal irreversibility threshold. That threshold is domain-specific and should be specified for the relevant decision context rather than inferred after the fact.

6 Independence, Redundancy, and Resilience

A central distinction is

$$\text{redundancy} \neq \text{independence}.$$

Multiple reviewers, control layers, or oversight channels do not establish resilience merely by being numerous. If the conditions that support Human NO remain concentrated in a single relevant common cause, multiple nominal layers may fail together.

RET asks whether refusal is effective *now*. Resilience asks whether the conditions that make RET true are excessively concentrated in one relevant failure source. Thus,

$$\boxed{\text{more layers} \not\Rightarrow \text{more resilience}}.$$

Independence is not absolute. No socio-technical architecture is independent of every physical, institutional, computational, or informational dependency.

The narrower public claim is:

Effective Human standing should not, where avoidable, depend on a single relevant common cause whose loss automatically erases all meaningful refusal capacity.

This paper does not publish system-specific dependency maps, exploitation paths, or bypass procedures. Those details belong to restricted validation practice, not to the public theory.

7 Dynamic RET

RET is not a static certification:

$$RET(S, t_0, \Omega) = TRUE$$

does not imply

$$RET(S, t_1, \Omega) = TRUE.$$

Nor does a result automatically transfer from one materially changed system or operating scope to another. Effective refusal may change because of changes in system capability, autonomy, authority scope, infrastructure, human organization, information dependencies, or operating environment.

The central assessment and governance sequence is

$$\boxed{\text{Human standing} \rightarrow RET(S, t, \Omega) \rightarrow \text{resilience of effective refusal} \rightarrow \text{revalidation} / \text{re-authorization}}$$

where the arrows represent sequence, not logical implication. Human standing does not imply a RET pass; a RET pass does not imply permanent resilience; and resilience does not eliminate the need for later revalidation.

8 Re-Authorization

Re-authorization is a temporal checkpoint. Its purpose is to prevent a past authorization or a past RET result from being silently carried forward after the conditions supporting it have materially changed. The governing rule is

$$\boxed{\text{material change} \Rightarrow \text{revalidation} / \text{re-authorization}}.$$

Re-authorization is not ceremonial approval. It cannot cure a failed RET by declaration alone. If effective refusal is no longer demonstrated, relevant deficiencies must first be addressed and RET revalidated before re-authorization can carry operational meaning.

This structure is consistent with governance approaches that treat monitoring, recovery, change management, contestability, and temporal renewal as continuing processes rather than one-time events [4, 10].

For future generations, the framework makes only a limited claim. Present institutions cannot literally obtain future consent. They can, however, avoid unnecessarily foreclosing the future capacity to contest and re-authorize persistent authority structures. UNESCO's 1997 Declaration connects responsibilities to future generations with preserving their freedom of choice over political, economic, and social systems [8].

9 Human Standing Under Capability Asymmetry

The framework does not claim that human judgment is always epistemically superior. As AI capability rises, human error-correction value may decline in some domains. The paper’s narrower normative distinction is

$$\boxed{\text{epistemic superiority} \neq \text{constitutive authority}}.$$

Being better at prediction, planning, or optimization does not by itself establish unilateral authority to erase the standing of affected humans.

Human NO is not an absolute veto. Emergencies, incapacity, serious third-party harm, and competing rights may justify override under separately specified legal or ethical standards. At broader social or civilizational scale, the paper does not define a single actor who “speaks for humanity.” The framework remains a test of effective refusal and standing, not a complete constitutional theory of AI governance.

10 Civilizational-Scale Governance Extension

At broader social or civilizational scale, the framework does not introduce a second operational test. It retains only the following governance considerations:

- broader and more irreversible effects increase concern about unilateral authorization;
- Human standing may need to be represented plurally rather than through one actor;
- sufficient information and time for refusal remain necessary;
- intervention must be effective rather than symbolic;
- prior refusal should not be rendered meaningless by automatic continuation after material change; and
- future re-authorization capacity should remain meaningful.

The exact institutional design, quorum, credential structure, or representative mechanism is outside this paper’s core contribution.

11 Practical Use

RET can be used as a design-review, assurance, audit, procurement, deployment, and change-management framework. A practical assessment should identify the specific system or version S , evaluation time t , operating scope Ω , evidence for each gate, known limitations and unknowns, and conditions that would require revalidation.

A RET result should remain scoped. A previous demonstration of effective refusal should not be treated as a permanent certificate after material changes. This makes RET compatible with lifecycle-oriented approaches such as the NIST AI RMF [4]. RET can also complement regulatory human-oversight requirements under Article 14 of the EU AI Act [5]: the additional operational question is whether that authority is actually usable in time and capable of changing the consequential trajectory.

12 Limitations

The framework does not imply unanimous global consent, permanent minority veto, world government, human infallibility, prohibition of autonomous AI action, future-generation interests overriding all present rights, or shutdown as the only refusal action.

It also does not claim absolute invulnerability against superintelligence, a theorem for containing arbitrarily capable systems, permanent Human supremacy, absolute independence from

every common cause, one-time RET certification as permanent validity, or a mathematical proof of normative Human standing.

The framework remains sensitive to emergencies, incapacity, public-safety externalities, conflicting human rights, latency constraints, refusal abuse, institutional capture, and domains in which maintaining refusal itself creates substantial costs or risks.

Most importantly:

The framework does not guarantee Human NO against an arbitrarily capable system. It tests whether Human standing and effective refusal are structured so that their loss is not made easy, silent, single-point, and irreversible.

13 Public-Safety Boundary

RET has a genuine dual-use dimension. The same categories that help defenders examine whether refusal is real can, if operationalized too specifically in public, reveal how superficial compliance might be engineered around an evaluation.

The public paper therefore publishes the seven defensive properties, timing and causal-efficacy principles, the requirement to assess relevant common-cause dependence, the dynamic nature of RET, the need for revalidation and re-authorization, and evidence obligations and limitations.

It does not publish exploitation-ready bypass procedures, exact failure-inducing sequences, system-specific critical weak points, detailed timing-manipulation procedures, concrete methods for preserving formal approval while defeating causal efficacy, or attack-oriented dependency maps.

The publication rule is

publish required properties and evidence obligations

while sequestering operational details that would materially increase evasion or bypass capability. This manuscript does not incorporate non-public material from the separate fourth-paper line of research.

14 Discussion

The central distinction is

human presence $\neq$ effective Human NO.

Existing MHC, oversight, and contestability research provides a strong foundation for human control, intervention, and accountability. Existing regulation already requires meaningful forms of human oversight and challenge. The Refusal Option framework narrows the focus to the moment before a consequential transition becomes difficult or impossible to reverse.

That focus produces four practical consequences. First, time is not an implementation detail; it is part of control itself. Second, causal efficacy matters more than nominal authorization. Third, redundancy must not be mistaken for independence. Fourth, validation expires when the system and its environment materially change.

The framework is intentionally modest about post-threshold guarantees. It does not claim that humans can permanently constrain an arbitrarily capable system. Its aim is to identify and audit conditions under which Human standing does not disappear automatically as capability and autonomy increase.

15 Conclusion

The operational core of the paper is:

Within the remaining time before irreversibility, can an independent Human NO actually change the outcome?

RET answers that question through seven non-compensatory gates:

$$\boxed{RET(S, t, \Omega) = G_1 \wedge G_2 \wedge G_3 \wedge G_4 \wedge G_5 \wedge G_6 \wedge G_7}.$$

The result is neither a safety certificate nor a proof of control over superintelligence. It is a scoped test of effective refusal.

The broader framework then asks whether that effectiveness is resilient to relevant common causes, whether changes are revalidated, and whether human standing can be re-authorized rather than silently erased.

The paper's strongest defensible claim is therefore limited:

The Refusal Option is a framework for examining whether the loss of effective Human standing can be prevented from becoming easy, silent, single-point, and irreversible before consequential AI-enabled transitions pass beyond meaningful human intervention.

Project Resources

The archived version of this paper is available at the Zenodo DOI below.

<https://doi.org/10.5281/zenodo.22980039>

LUMINA-30 project: <https://github.com/lumina-30/lumina-30-overview>

PDF archive: <https://github.com/lumina-30/Lumi30-PDF-Archive>

Related pre-irreversibility research (PCR-C):

<https://doi.org/10.5281/zenodo.18824181>

This paper is provided as a public-safe research framework. It does not create certification status, compliance obligations, deployment authorization, or institutional approval under LUMINA-30.

References

- [1] F. Santoni de Sio and J. van den Hoven. Meaningful Human Control over Autonomous Systems: A Philosophical Account. *Frontiers in Robotics and AI*, 5:15, 2018. <https://doi.org/10.3389/frobt.2018.00015>.
- [2] G. Mecacci and F. Santoni de Sio. Meaningful human control as reason-responsiveness: the case of dual-mode vehicles. *Ethics and Information Technology*, 22:103–115, 2020. <https://doi.org/10.1007/s10676-019-09519-w>.
- [3] H. Veluwenkamp. Reasons for Meaningful Human Control. *Ethics and Information Technology*, 24:51, 2022. <https://doi.org/10.1007/s10676-022-09673-8>.
- [4] E. Tabassi. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1, National Institute of Standards and Technology, 2023. <https://doi.org/10.6028/NIST.AI.100-1>.

- [5] European Union. Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act), Article 14, 2024. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>.
- [6] Council of Europe. *Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law*. CETS No. 225, 2024. <https://www.coe.int/en/web/artificial-intelligence/the-framework-convention-on-artificial-intelligence>.
- [7] UNESCO. *Recommendation on the Ethics of Artificial Intelligence*. Adopted 23 November 2021. <https://www.unesco.org/en/artificial-intelligence/recommendation-ethics>.
- [8] UNESCO. *Declaration on the Responsibilities of the Present Generations Towards Future Generations*. 1997. <https://www.unesco.org/en/legal-affairs/declaration-responsibilities-present-generations-towards-future-generations>.
- [9] J. A. Torrecilla-Pinero. Judgment Cannot Be Delegated: A Subject-Preserving Framework for AI Governance. *Philosophy & Technology*, 39:135, 2026. <https://doi.org/10.1007/s13347-026-01140-2>.
- [10] G. Abukar. Temporal authorship and the moral illegitimacy of pre-authorized AI: a new ethical architecture for delegated inquiry. *AI and Ethics*, 6:215, 2026. <https://doi.org/10.1007/s43681-026-01080-9>.
- [11] D. van de Sande, N. Economou-Zavlanos, and M. E. van Genderen. Meaningful oversight of medical AI beyond human in the loop. *npj Digital Medicine*, 9:569, 2026. <https://doi.org/10.1038/s41746-026-02971-1>.
- [12] L. Zhu, Q. Lu, M. Ding, S. U. Lee, et al. Designing meaningful human oversight in AI. *AI and Ethics*, 6:286, 2026. <https://doi.org/10.1007/s43681-026-01147-7>.